

Review of Material from Lecture 2

EL GY 6463: Advanced Hardware Design

Instructor: Siddharth Garg

Finite State Machines (Data path and Control path Synthesis)


- Say you want to build a multiplier but only have access to adders and comparators.
 - Solution: repeated addition. Simple Verilog combinational description

```
module mult #(parameter W = 8)
    (input [W-1:0] x, input [W-1:0] y; output reg [2*W-1:0] out);

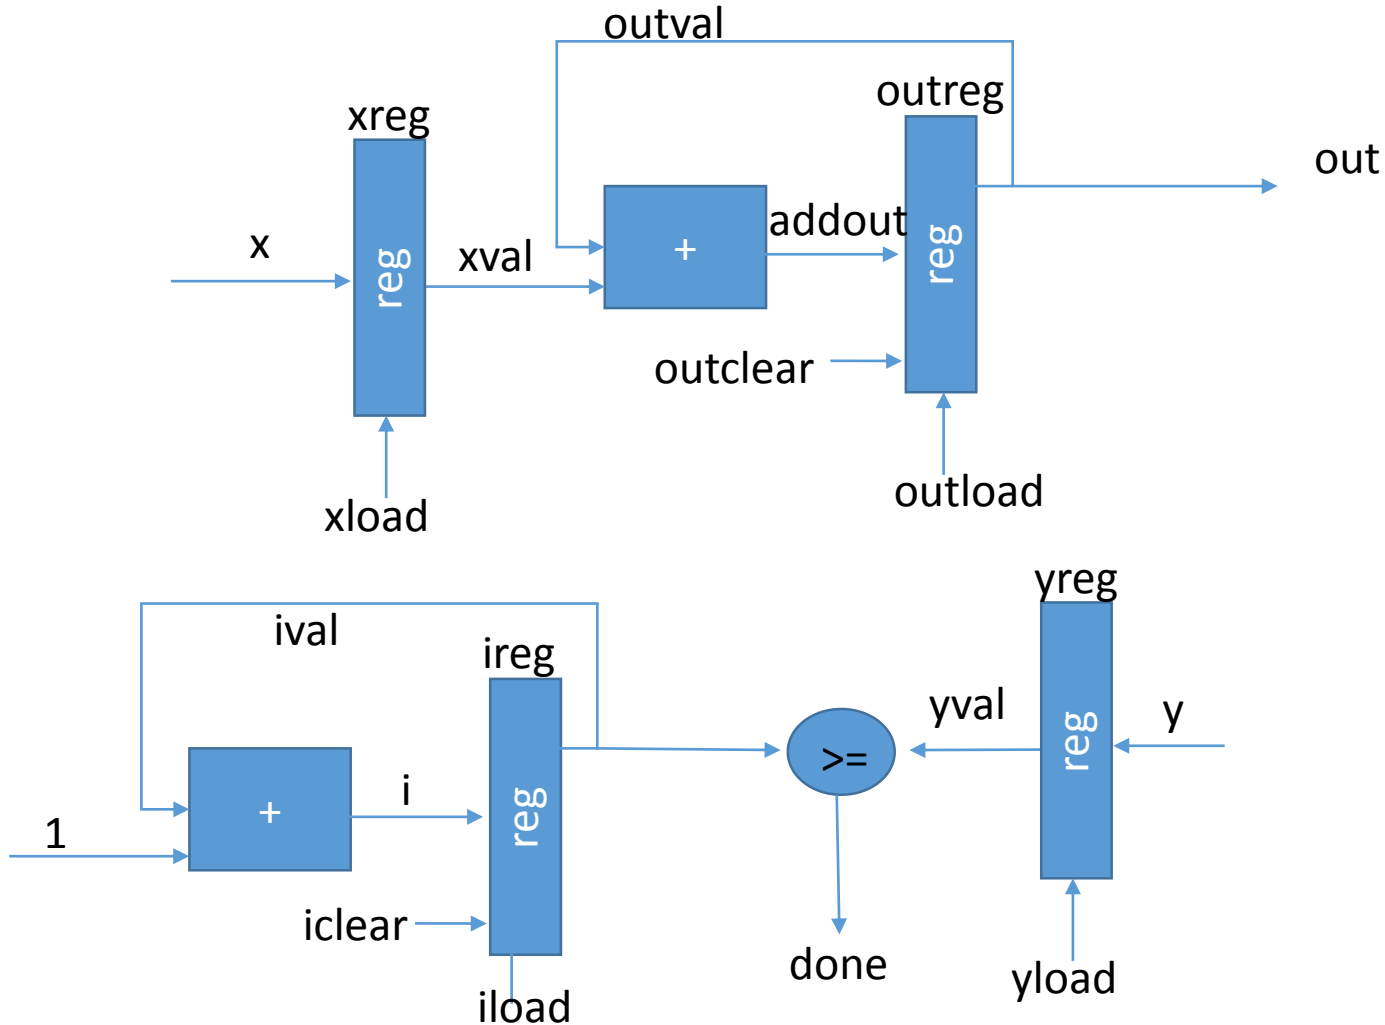
    reg [W-1:0] i;    Will work in
                     simulation but does
                     not synthesize?

    always @(*)
        out = 0;
        for (i=0; i<y; i=i+1) begin
            out = out + x;
        end

endmodule
```



Better Solution: Use One Adder Repeatedly



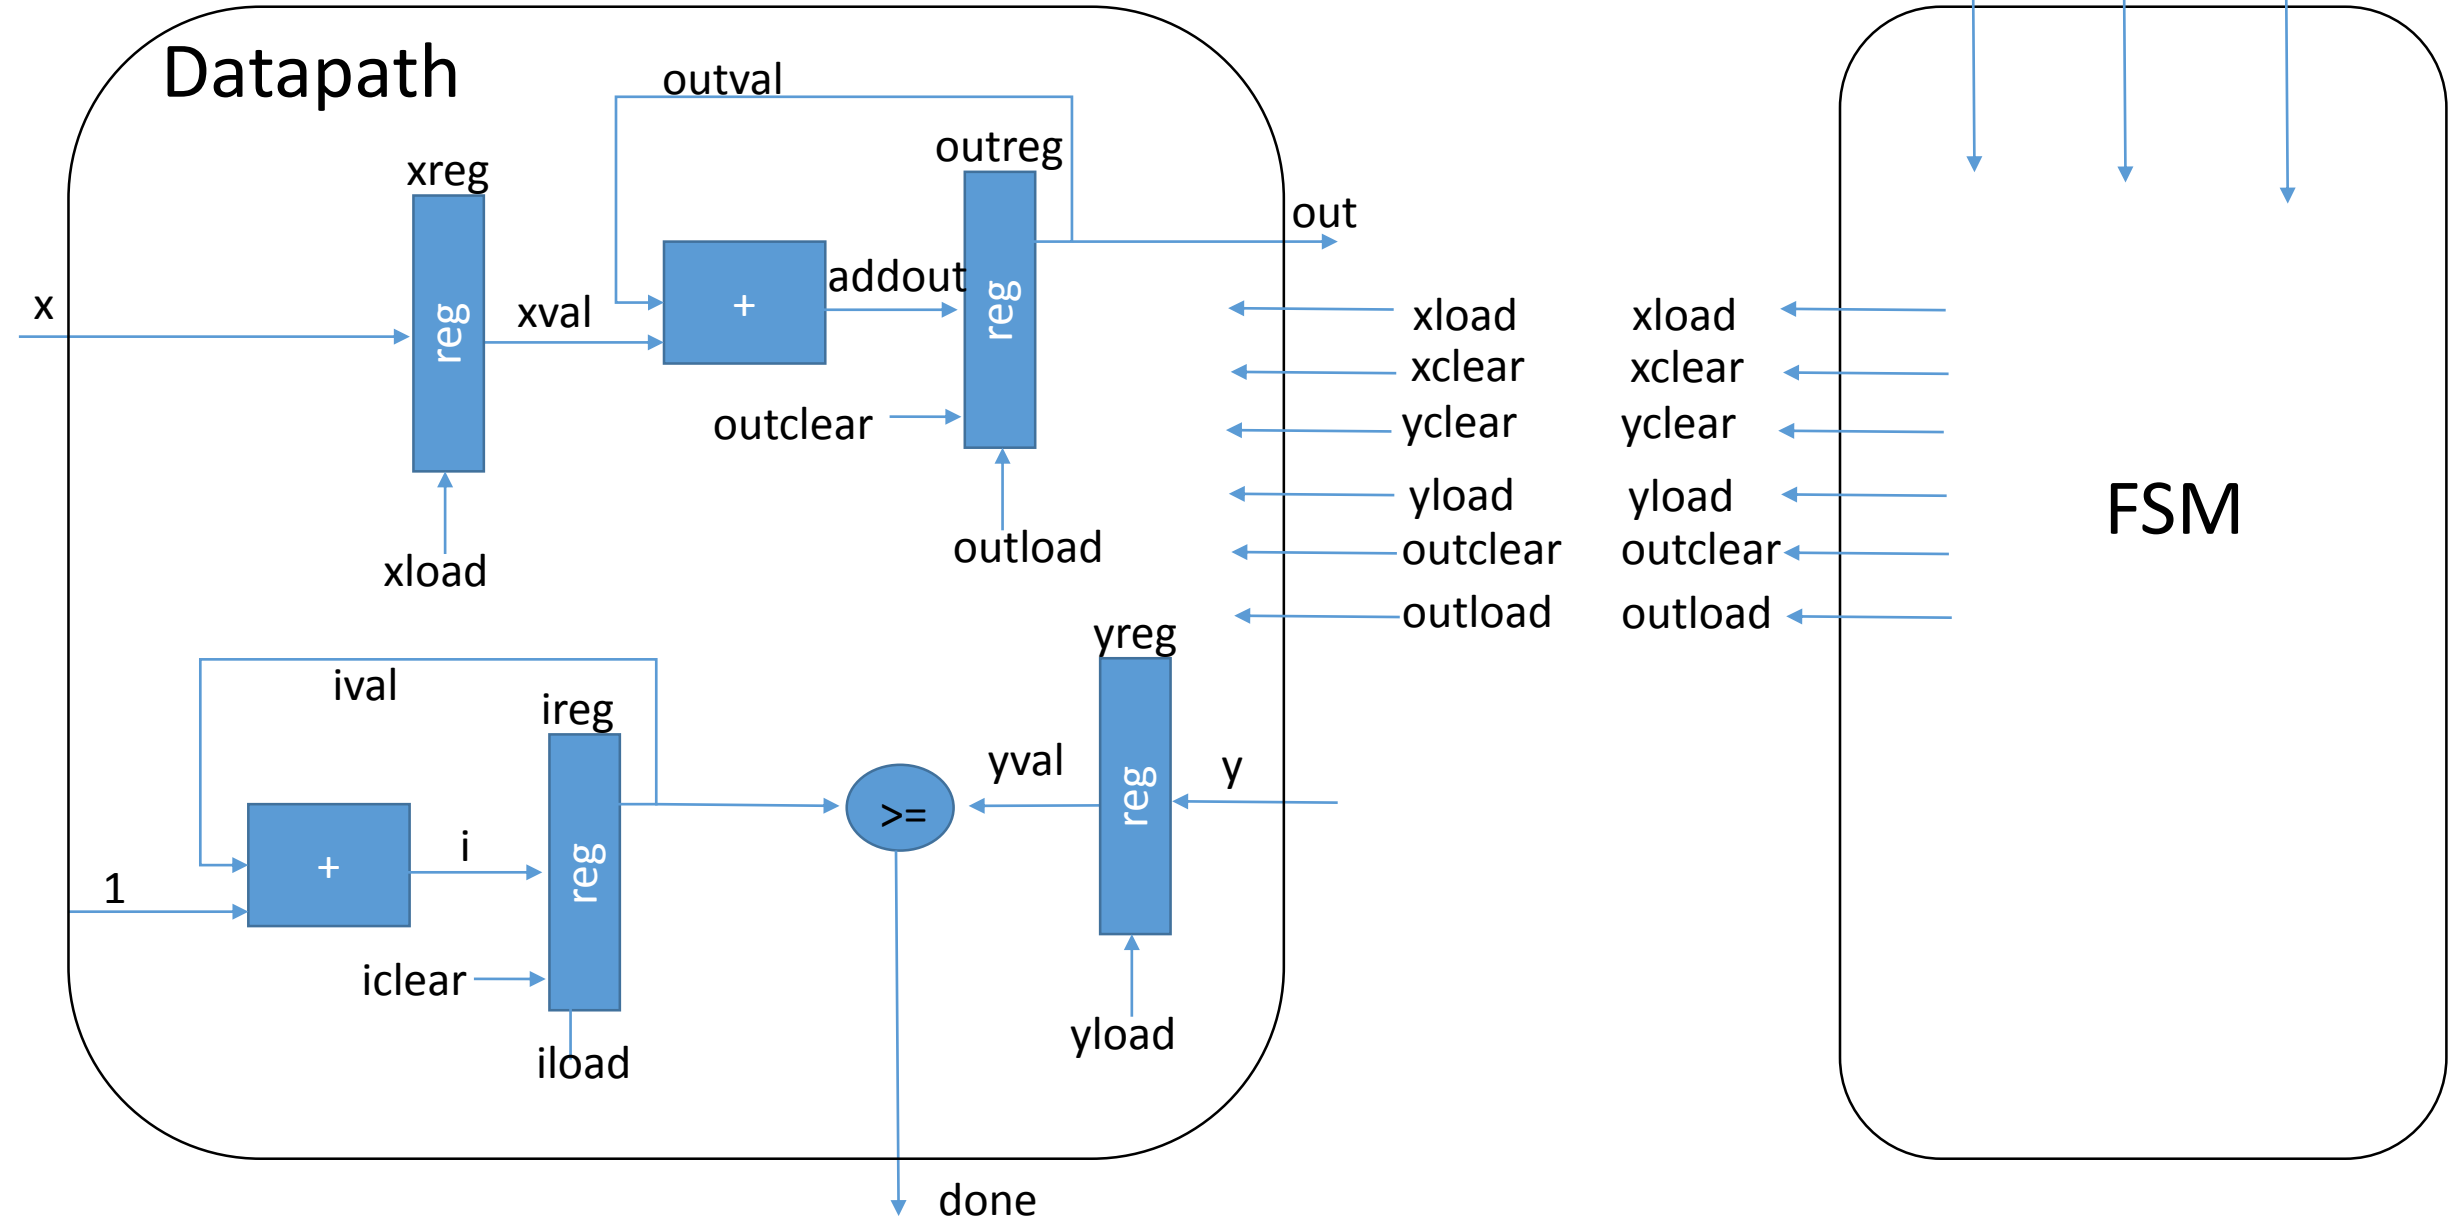
```
module increment
# (parameter W = 8)
(input [W-1:0] in1,
output reg [W-1:0] out);
```

```
always @ (*) begin
    out = in1 + 1;
end
endmodule
```

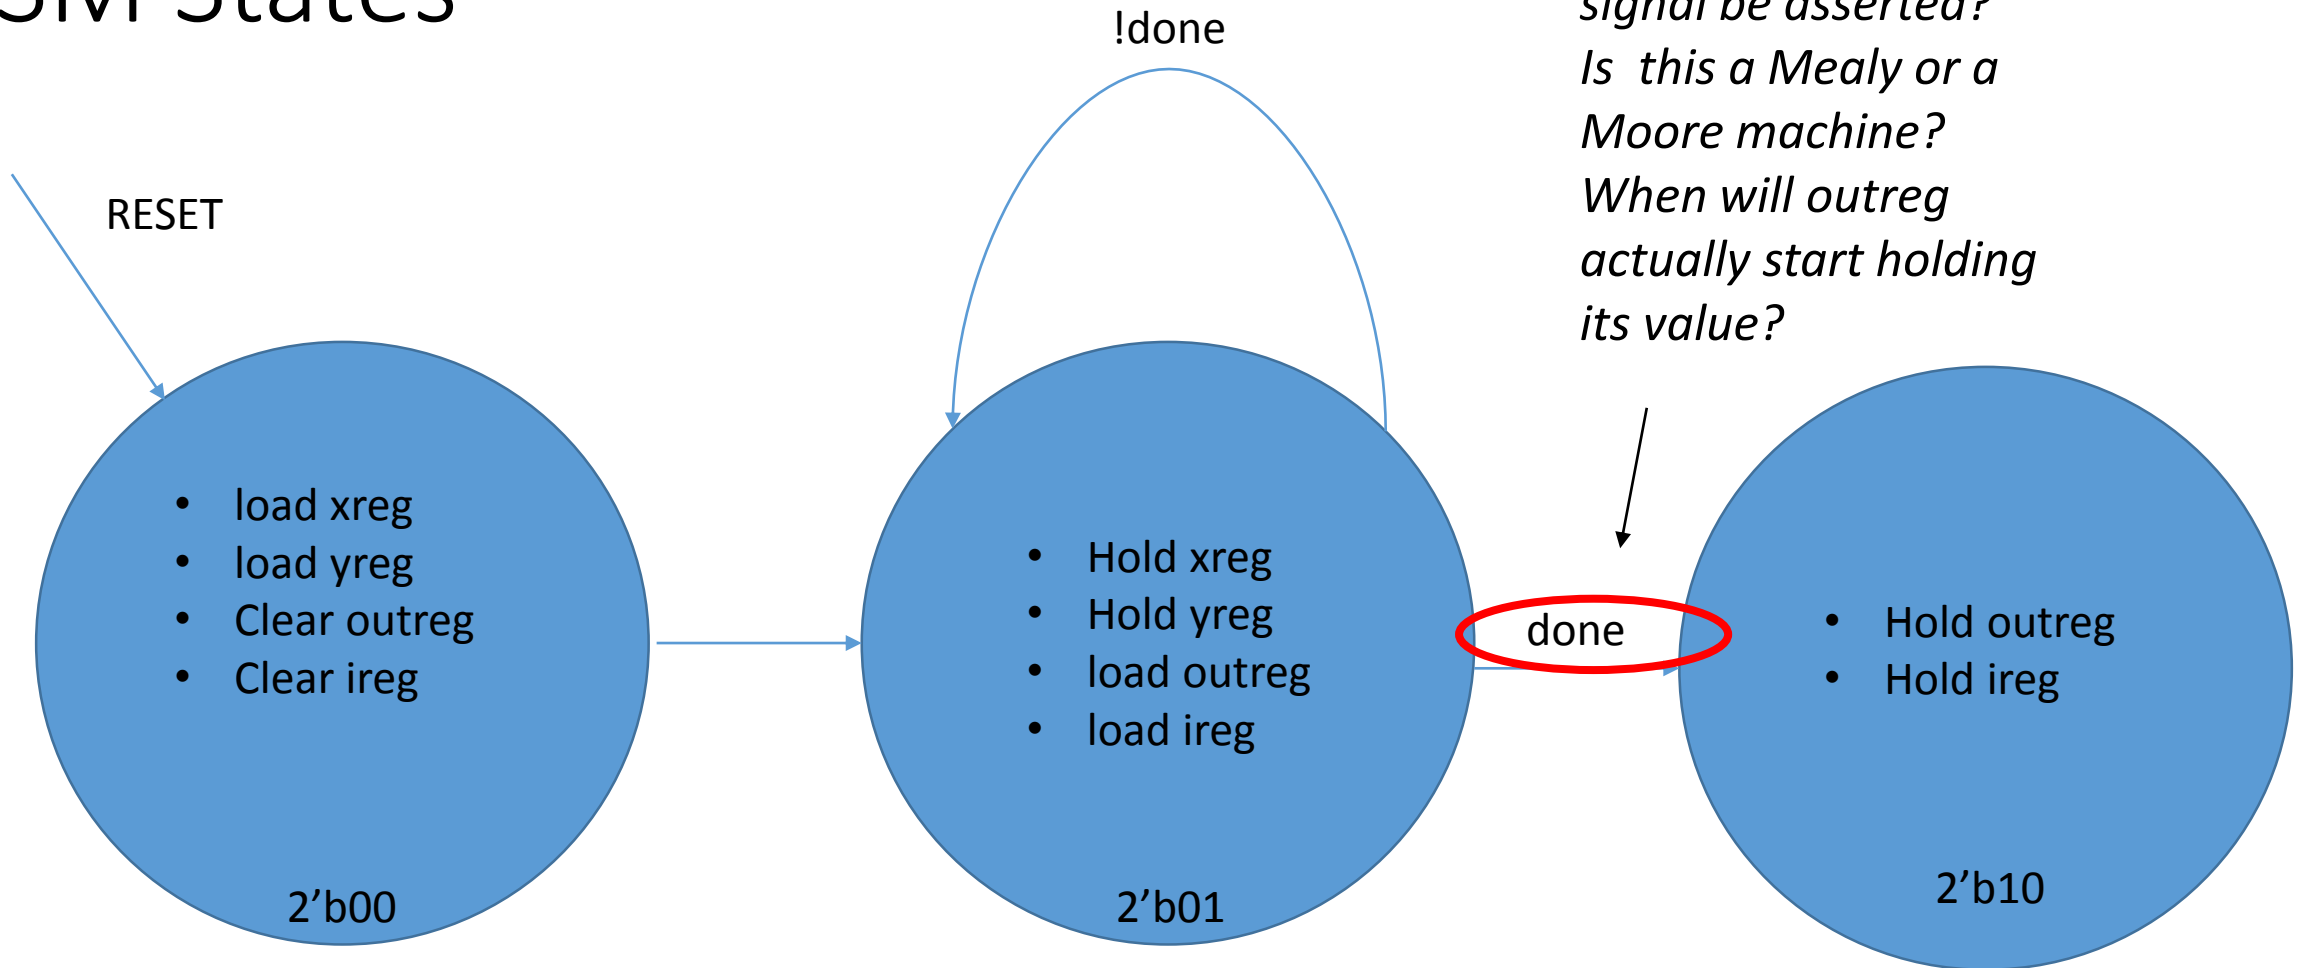
```
module comparator
# (parameter W = 8)
(input [W-1:0] ival,
input [W-1:0] yval,
output reg done);
```

```
always @ (*) begin
    done = 0;
    if (ival >= yval ) //Yes?
        done=1;
end
endmodule
```

How to Control the Datapath?



FSM States



Increment Module

```
module comparator
# (parameter W = 8)
(input [W-1:0] ival,
 input [W-1:0] yval,
 output reg done);

always @ (*) begin
    done = 0;
    if (ival +1    >= yval -1 )
        done=1;
end
endmodule
```

Note:

- ival starts from 0 instead of 1
- When done signal is asserted, outreg only stops loading in the next clock cycle

Putting it All Together

```
module multiplier
# (parameter W = 8)
(input [W-1:0] x, input [W-1:0] y, input clk, input reset, output [2*W-1:0] out);

//variable declarations. Question: what datatype for xval, xclear, xload?

assign out = addout;

myreg #(W) xreg(x, xload, xclear, clk, xval); //xreg
myreg #(W) yreg(y, yload, yclear, clk, yval); //yreg
myreg #(2*W) outreg(addout, outload, outclear, clk, outval); //outreg
myreg #(W) ireg(i, iload, iclear, clk, ival); //ireg

adder #(W) add1 (xval, outval, addout); //out = out+x
increment incl (ival, i); //i = i+1
comparator #(W) comp1 (ival, yval, done); //i+1 >= y-1

myfsm fsm1 ( clk, reset, done, xload, xclear, yload, yclear, iload, iclear,
outload, outclear); //controller

endmodule
```

FSM Design (In-Class Exercise)

```
module myfsm (input clk, input reset, input done, output reg xload, output...);

reg [2:0] cs, ns;

always @ (posedge clk, posedge reset) begin
    //SEQUENTIAL PORTION: YOUR CODE HERE
end

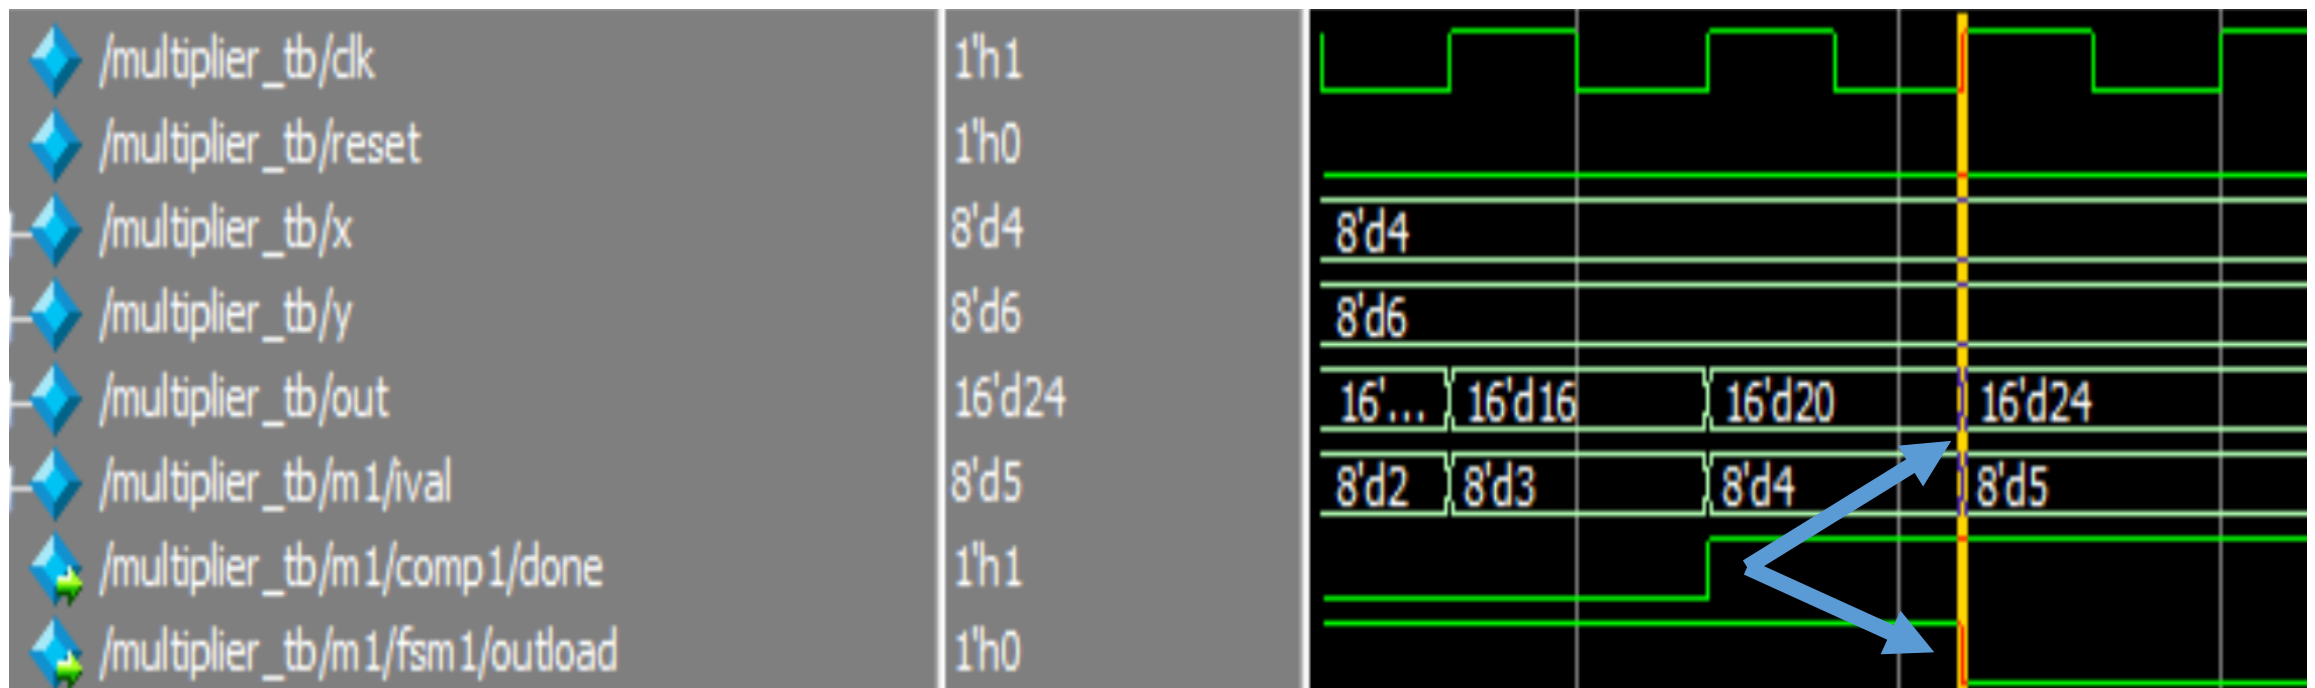
always @(*) begin
    ns = cs;
    xload = 0; yload = 0; iload = 0; outload = 0;
    xclear = 0; yclear = 0; iclear = 0; outclear = 0;
    case (cs)
        2'b00 : begin
            // YOUR CODE HERE
        end
        2'b01: begin
            //YOUR CODE HERE
        end
        2'b10 : begin
            // YOUR CODE HERE
        end
        default :
    endcase
end
endmodule
```


Solution

```
case (cs)
    //START
    2'b00 : begin
        xload = 1; yload = 1;  iclear = 1; outclear = 1;
        ns = 2'b01;
    end
    //working
    2'b01: begin
        xload = 0; yload = 0; iload = 1; outload = 1;
        if (done)
            ns = 2'b10;
        end
    //finished
    2'b10 : begin
        iload = 0; outload = 0;
    end

    default : // put in don't cares
```

Desired Simulation Output



One clock cycle difference. Think of done as “almost done”

Assignment 1: Review

Problem 1

- Huffman Decoder. What is a Huffman code?

A naive code would use the same number of bits to represent each symbol. For instance, if we have 4 symbols {A,B,C,D}, a naive code would use 2 bits for each.

{A = 00, B = 01, C = 10, D = 11}.

But what if some symbols are more frequent than others?

AAABAAAADACAAABBBAAA

Idea: Use fewer bits for the most frequent symbols.

Example Codebook

Codeword	Original Symbol
0	A
10	B
110	C
111	D

Let's say symbol A occurs with 70% likelihood and all other symbols occur with 10% likelihood. **What is the compression ratio compared to a Naïve encoding?**

Prefix Code. A Huffman code is an example of a Prefix code. That is, no codeword is the prefix of another codeword.


Example, let's look at codeword 110 (C). Prefixes of this codeword:

- 1 (not a codeword)
- 11 (not a codeword)
- 110 (Codeword = C)

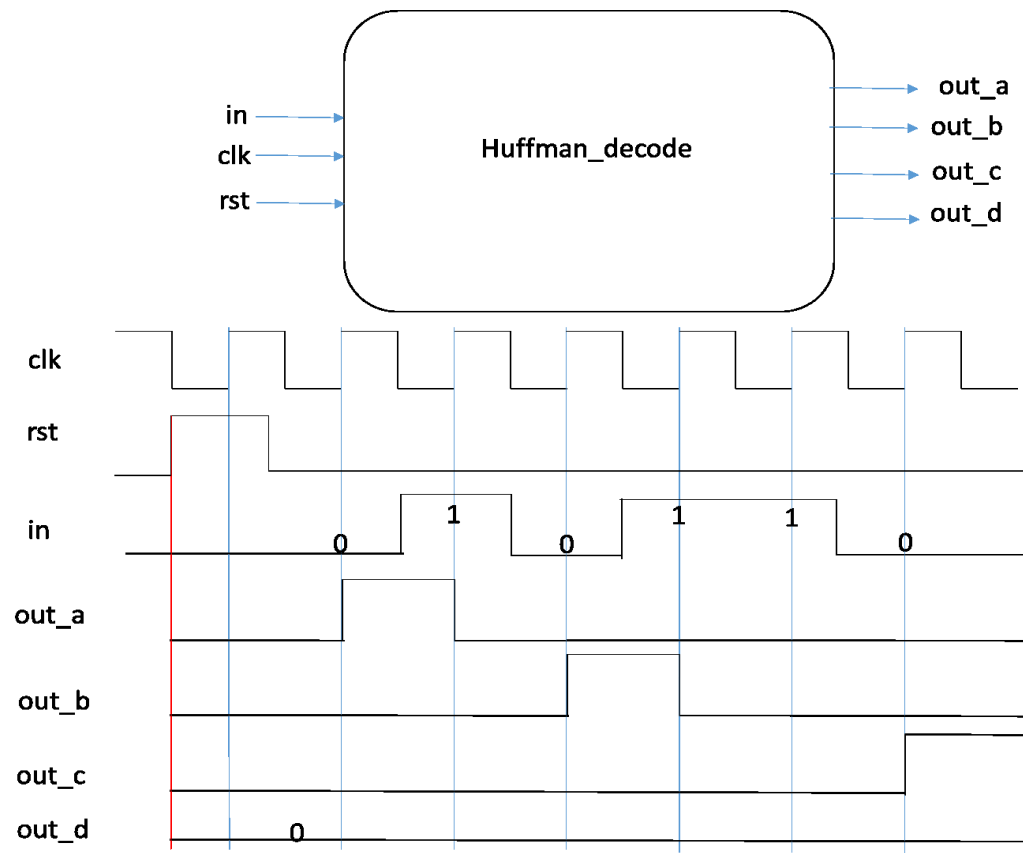
This means symbols can be decoded as bits from the compressed stream arrive! No need to know what the future bits will be.

- Consider what would happen if the codeword for A were 1!

0 0 1 1 0 1 0 1 1 1 0
A A - - C - B - - D A



Understanding Specifications



Question: is this a Mealy or a Moore machine?
(when do the outputs change?)

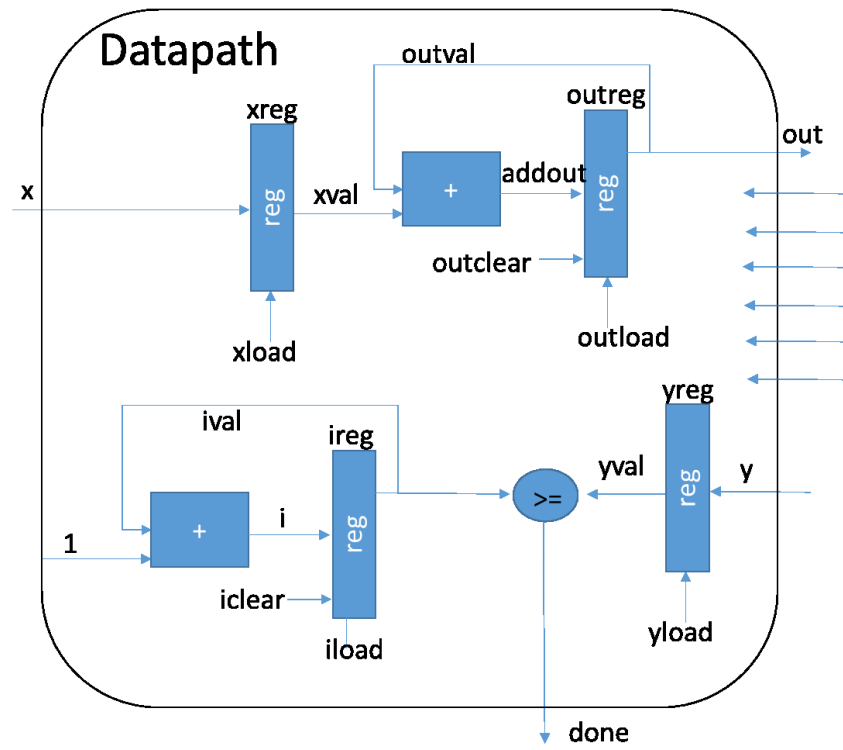
Problem 2: Divider

```
out = 0;  
for (i=0; i<y; i=i+1) begin  
    out = out + x;  
end
```

Multiplier Code

```
quotient = 0;  
while (a>=b) {  
    a = a - b;  
    quotient = quotient + 1;  
}  
remainder = a;
```

Divider Code



What will the datapath for the divider look like?

Lecture 3: Synthesis

EL GY 6463: Advanced Hardware Design

Instructor: Siddharth Garg

What is Synthesis?

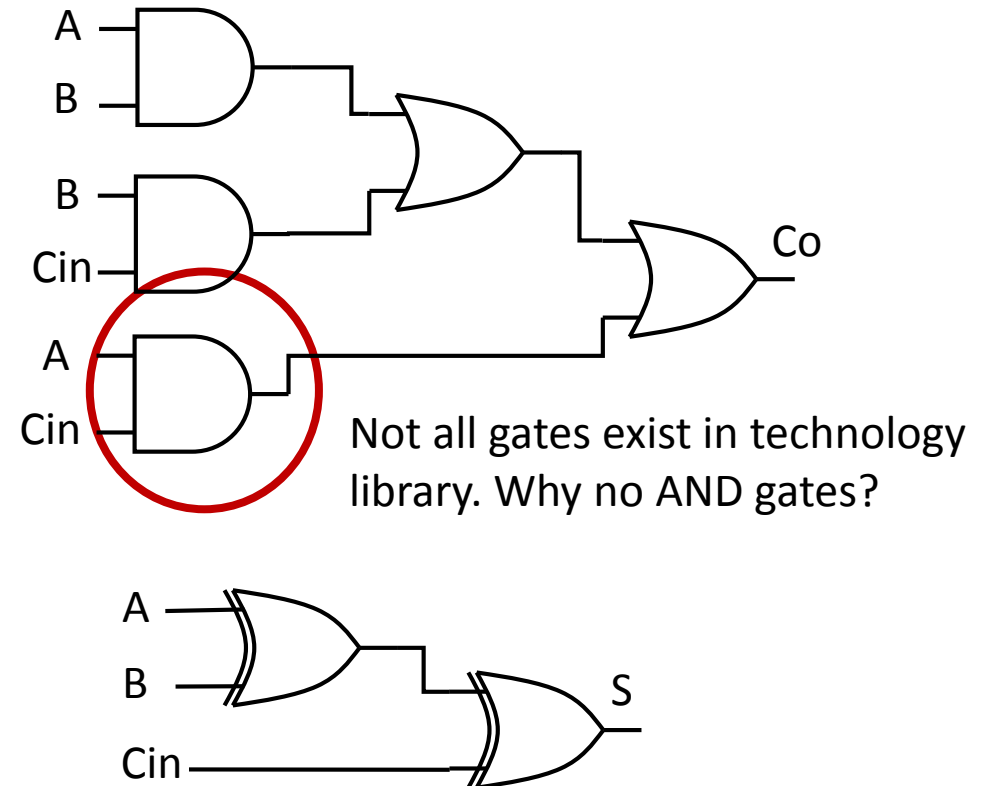

- Converts a Verilog description of logic functionality (combinational or FSM) to a netlist of logic Boolean gates + sequential elements like latches and flip-flops
 - The types of logic gates and latches/flip-flops is specified in a **technology library**

```
module fulladder ( input a,
                  input b,
                  input cin,
                  output sum,
                  output cout);

  assign {cout,sum} = a + b + cin;

endmodule
```

Synthesis



Cadence RC Synthesis Tool Demo

- You all have accounts on gauss.poly.edu
 - Please try connecting and logging in. If you cannot login, please send ASAP an email to Julio Rivera (jr847@nyu.edu) and cc me.
- FAQ: can I get Cadence RC on my desktop? No. commercial tool, thousands of dollars per license


Full Adder Synthesis Output

```
module fulladder ( input a,
                  input b,
                  input cin,
                  output sum,
                  output cout);



  assign {cout,sum} = a + b + cin;

endmodule
```

Synthesis



```
module fulladder(a, b, cin, sum, cout);
  input a, b, cin;
  output sum, cout;
  wire a, b, cin;
  wire sum, cout;
  wire n_0, n_1, n_2;
  XOR2X1 g56(.A (cin), .B (n_1), .Y (sum));
  NAND2X1 g57(.A (n_0), .B (n_2), .Y (cout));
  OAI21X1 g58(.A (a), .B (b), .C (cin), .Y (n_2));
  XOR2X1 g59(.A (a), .B (b), .Y (n_1));
  NAND2X1 g60(.A (b), .B (a), .Y (n_0));
endmodule
```

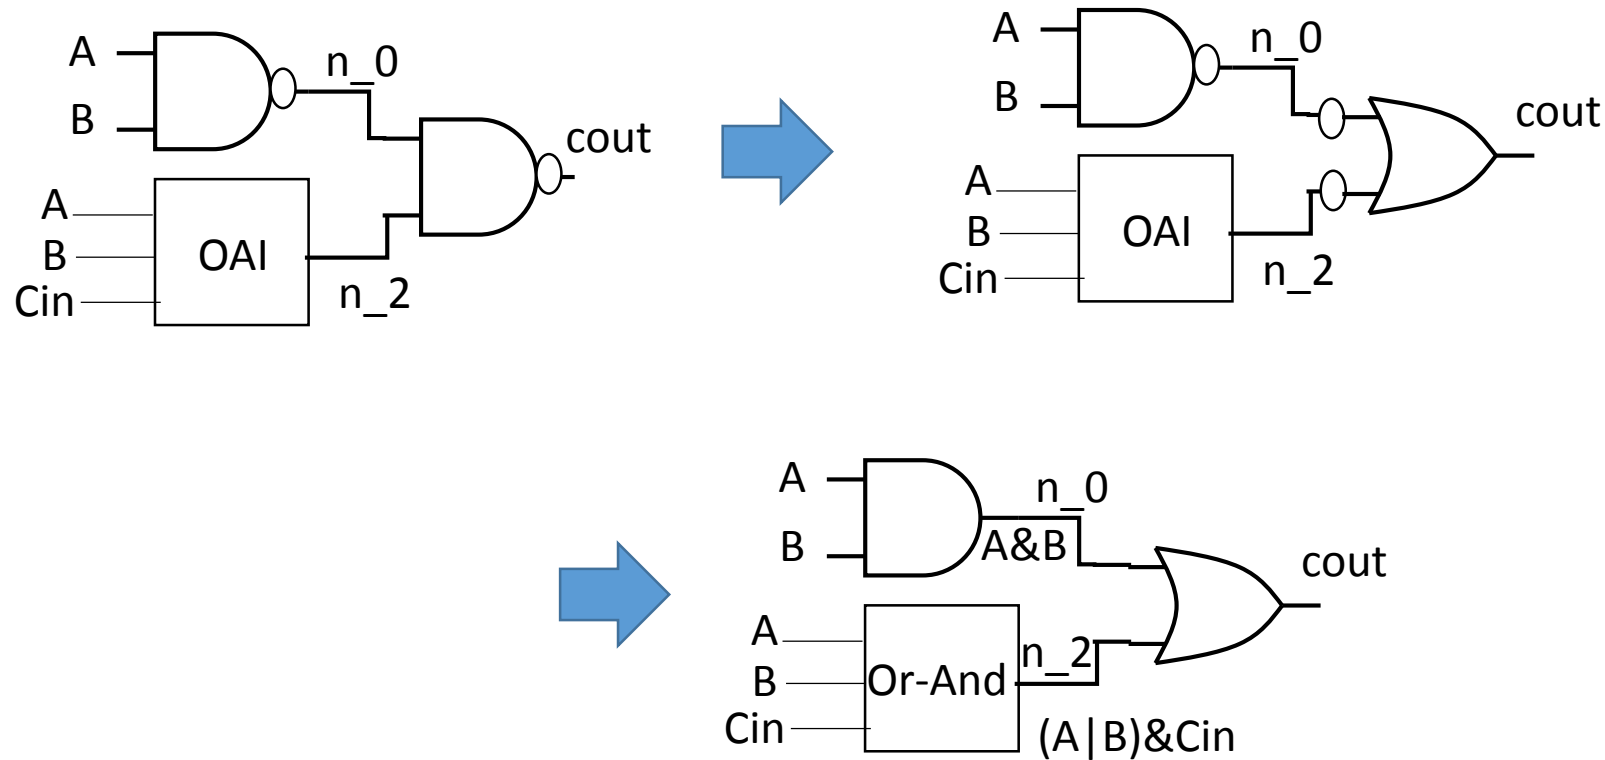


$$\text{OAI}(x,y,z) = ((x|y) \& z)'$$

In class exercise: prove that the synthesized netlist for cout gives the correct output. Recall that:
 $\text{cout} = a\&b \mid b\&c \mid c\&a$

What is an OAI gate?

Solution to In-class Exercise



Question: how did the synthesis tool choose this specific solution?

- What metrics is the tool trying to optimize for?

Optimization Metrics

- Three primary metrics that a synthesis tool cares about
 - **Area**: what is the total sum of areas of each gate
 - **Delay**: what is the delay from the input to the output (more on this very shortly)
 - **Power**: what is the circuit power consumption (full lecture dedicated to power consumption later in the term)
- Let's look at area first. How does the synthesis tool know the area of each gate?
 - Look in the **technology library**

Technology library

- RC needs us to tell it where the technology library exists


- set_attribute lib_search_path
- set_attribute library

- In “run.tcl”


```
set_attribute lib_search_path {/opt/cadence/local/ncsu-cdk-1.5.1-with-UofU/lib/UofU_Digital_v1_2/}
set_attribute library [list UofU_Digital_v1_2 modified.lib];
```

- Let's take a look at UofU_Digital_v1_2_modified.lib

```
/* ----- *
 * Design : OAI21X1 *
 * ----- */
(OAI21X1) {
  cell footprint : oai21;
  area : 5;
  cell_leakage_power : 0.0721417;
  pin(A) {
    direction : input;
    capacitance : 0.0308513;
    rise_capacitance : 0.0308513;
    fall_capacitance : 0.0306689;
    rise_capacitance_range ( 0.0287875, 0.0329151 ) ;
    fall_capacitance_range ( 0.0285731, 0.0327647 ) ;
    internal_power() {
      rise_power(passive_energy_template_5x1) {
        index_1 ("0.06, 0.18, 0.42, 0.6, 1.2");
        values ("0.482069, 0.481844, 0.481735, 0.481881,
      }
    }
  }
}
```



11



16

How Did the Synthesis Tool Get to This Solution?

- *Many, many* steps from a behavioral Verilog description to a structural description mapped to a given technology
 - Technology dependent **logic minimization** (this step does not depend on the technology library)
 - **Technology mapping** (only gates from the technology library allowed)
 - Further technology dependent optimizations like gate sizing (NANDx1, NANDx2...)


Two-level Logic Minimization

- Any Boolean function can be expressed in “*sum of products*” form
 - $F = (a \& b) \mid (b' \& c) \mid (a' \& b' \& c')$ (3 product terms, 7 literals)
 - Product term
 - Literal
- Goal of two level logic minimizations: **determine an SOP expression for F which has the *fewest product terms* OR *fewest literals***
 - Why should minimizing the number of product terms or literals matter?

Two-Level Minimization: Prime Implicants

- Start with all entries in the Truth table that output a '1' (ON-Terms)

W	X	Y	Z	Label
0	0	0	0	M0
0	1	0	1	M5
0	1	1	1	M7
1	0	0	0	M8
1	0	0	1	M9
1	0	1	0	M10
1	0	1	1	M11



W	X	Y	Z	Label
-	0	0	0	M0, M8
0	1	-	1	M5, M7
1	0	0	-	M8, M9
1	0	-	0	M8, M10
1	0	-	1	M9, M11
1	0	1	-	M10, M11

All pairs of terms that differ in a single entry

W	X	Y	Z	Label
1	0	-	-	M8, M9, M10, M11


All sets of 4 terms that differ in only two entries

Prime implicants are entries in these tables that are not subsets of other entries

Example modified from MIT OCW

Two-Level Minimization: Covering Problem

- For all prime implicants, list the min terms that they cover



PI	M0	M5	M7	M8	M9	M10	M11
{M0,M8}	1			1			
{M5, M7}		1	1				
{M8, M9, M10, M11}				1	1	1	1

- Every time you pick a PI, it “covers” all its minterms
 - Our goal: pick the smallest number of PIs so all minterms are covered.
 - How many in this example?

Solution

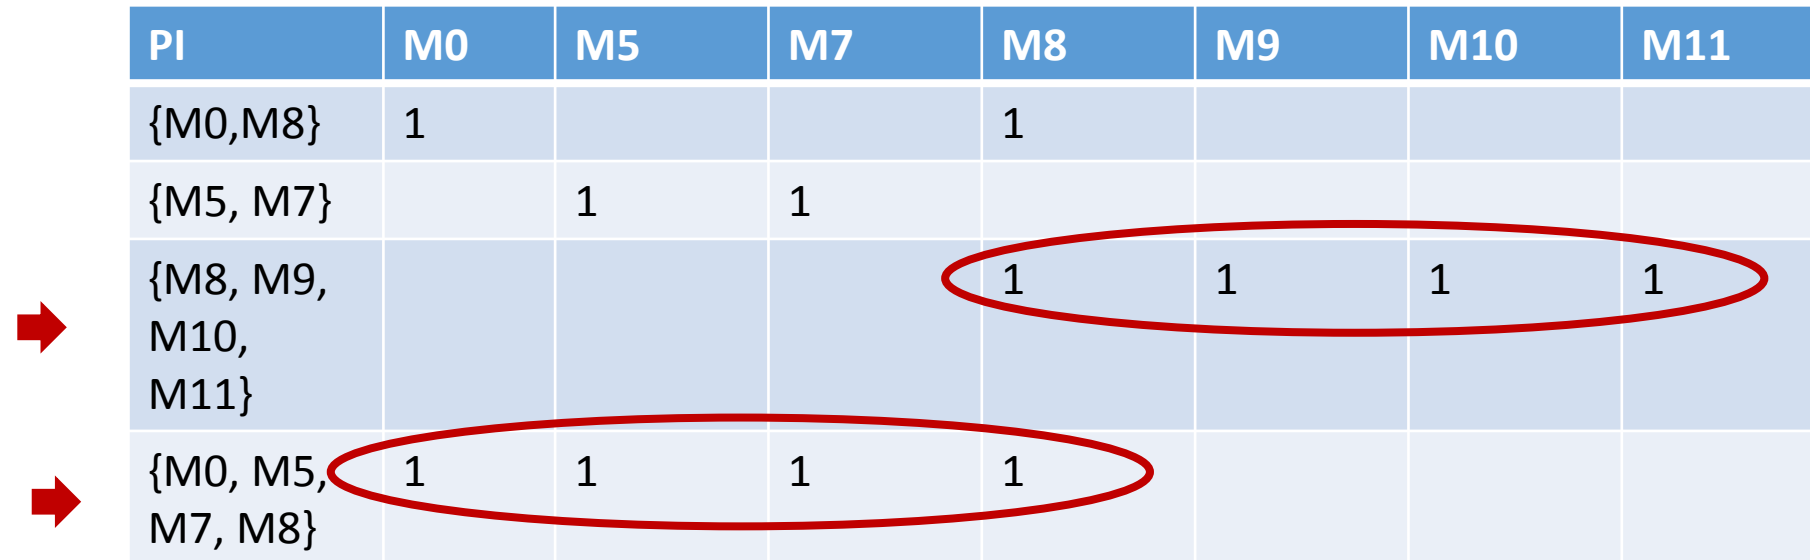
- In this example, all PIs need to be picked to cover all minterms

PI	M0	M5	M7	M8	M9	M10	M11
{M0,M8}	1			1			
{M5, M7}		1	1				
{M8, M9, M10, M11}				1	1	1	1

- $F = (x' \& y' \& z') \mid (w' \& x \& z) \mid (w \& x')$
 - How many two input AND/OR gates assuming true and complemented forms of inputs?

Set Cover

- But in general, not all min terms might be required for the cover





PI	M0	M5	M7	M8	M9	M10	M11
{M0,M8}	1			1			
{M5, M7}		1	1				
{M8, M9, M10, M11}				1	1	1	1
{M0, M5, M7, M8}	1	1	1	1			

- The set-cover problem in general is NP-complete
 - No poly-time solution unless P=NP
 - Cadence RC implements several heuristic solutions to solve such problems


Technology Mapping

- OK, so you now have a minimized SOP formula for your Boolean function.
 - How does that translate to gates in the technology library?
- **STEP 1:** represent your Boolean formula as a circuit consisting of “elemental” gates. Typically NAND and INVs
 - $F = (a \& b) \mid (b \& c) \mid (c \& a) = ((a \mid b) \& c) \mid (c \& a)$




Technology Mapping

STEP 2: represent every gate in your tech library with “elemental” gates. Typically NAND and INVs




STEP 3: match sub-graphs in F to gates in tech. library



Technology Mapping

STEP 2: represent every gate in your tech library with “elemental” gates. Typically NAND and INVs



STEP 3: match sub-graphs in F to gates in tech. library

○ INV (COST = 1)

NAND (COST = 3)


NOR (COST = 4)

NAND3 (COST = 6)

OAI3 (COST = 5)

Technology Mapping


STEP 2: represent every gate in your tech library with “elemental” gates. Typically NAND and INVs




Total cost = 4 + 6 + 3 = 13




INV (COST = 1)




NAND (COST = 3)



NOR (COST = 4)




NAND3 (COST = 6)



OAI3 (COST = 5)

Technology Mapping

STEP 2: represent every gate in your tech library with “elemental” gates. Typically NAND and INVs

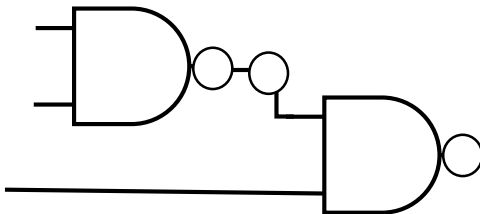


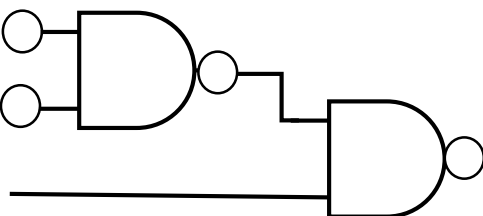
$$\text{Total cost} = 7 + 3 + 3 + 3 \times 1 = 16$$

○ INV (COST = 1)

 NAND (COST = 3)


 NOR (COST = 4)

 NAND3 (COST = 6)

 OAI3 (COST = 5)

Dynamic Programming

Is there a systematic way to find optimal mapping? Yes! (Dynamic programming)



○ INV (COST = 1)

AND NAND (COST = 3)

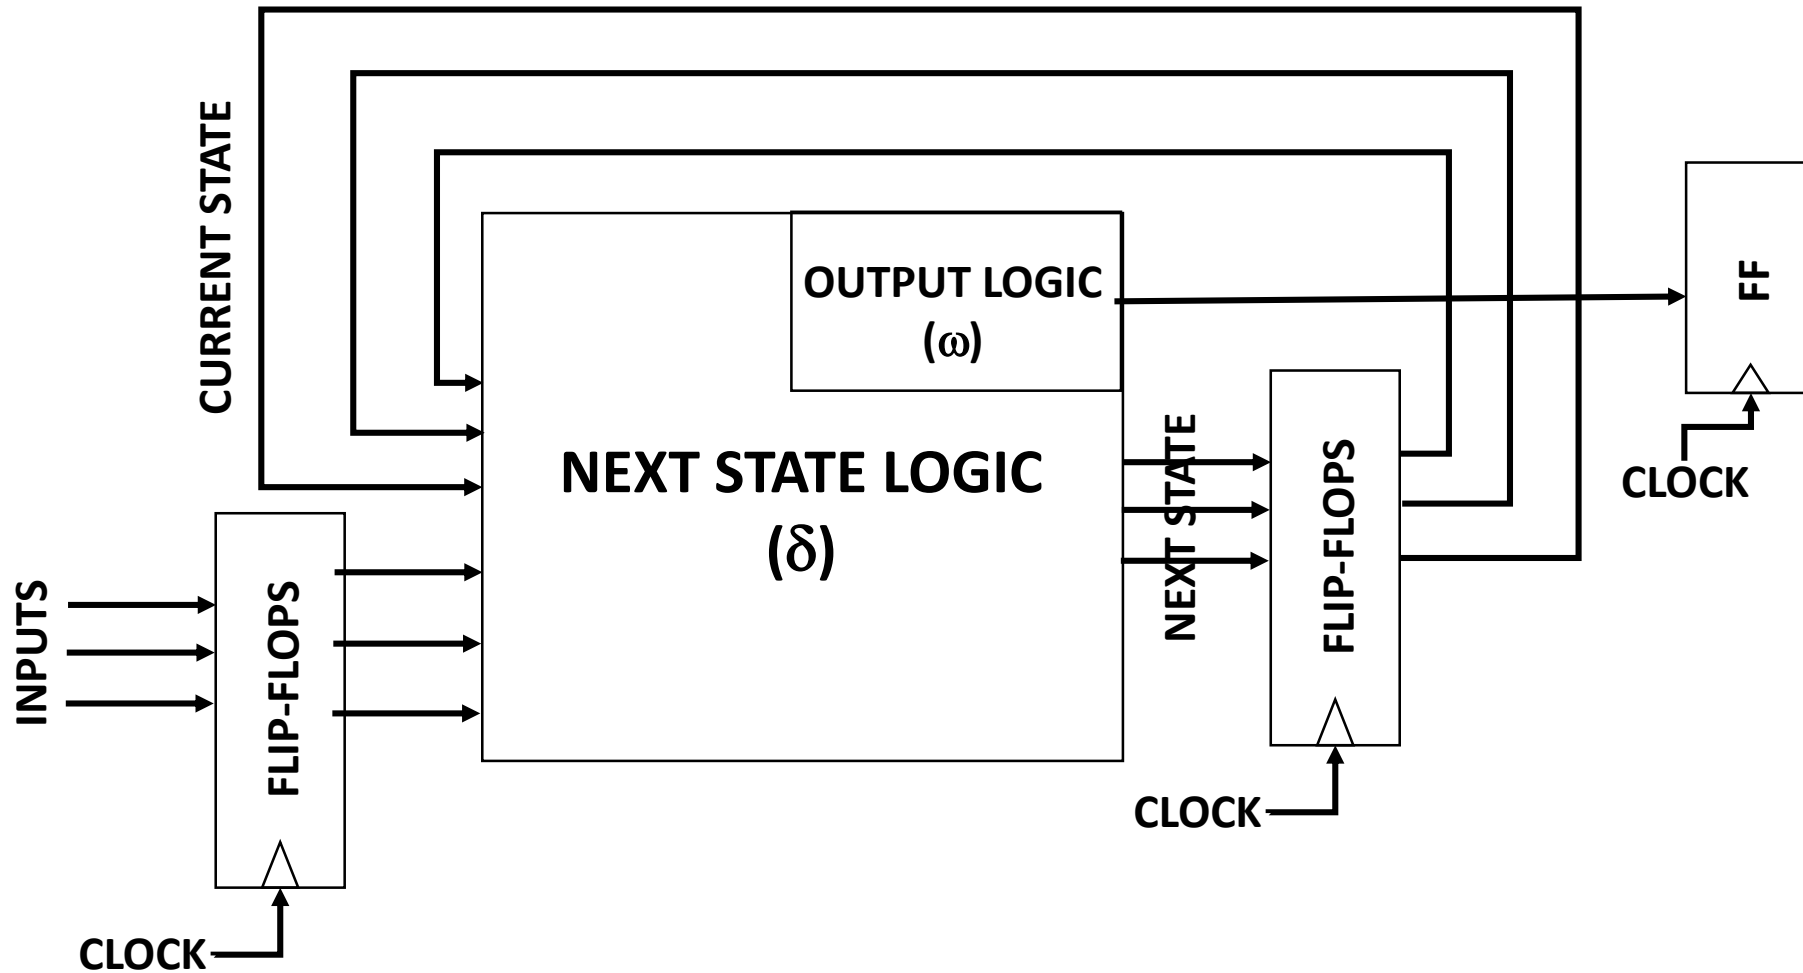
NOR (COST = 4)

NAND3 (COST = 6)

OAI3 (COST = 5)

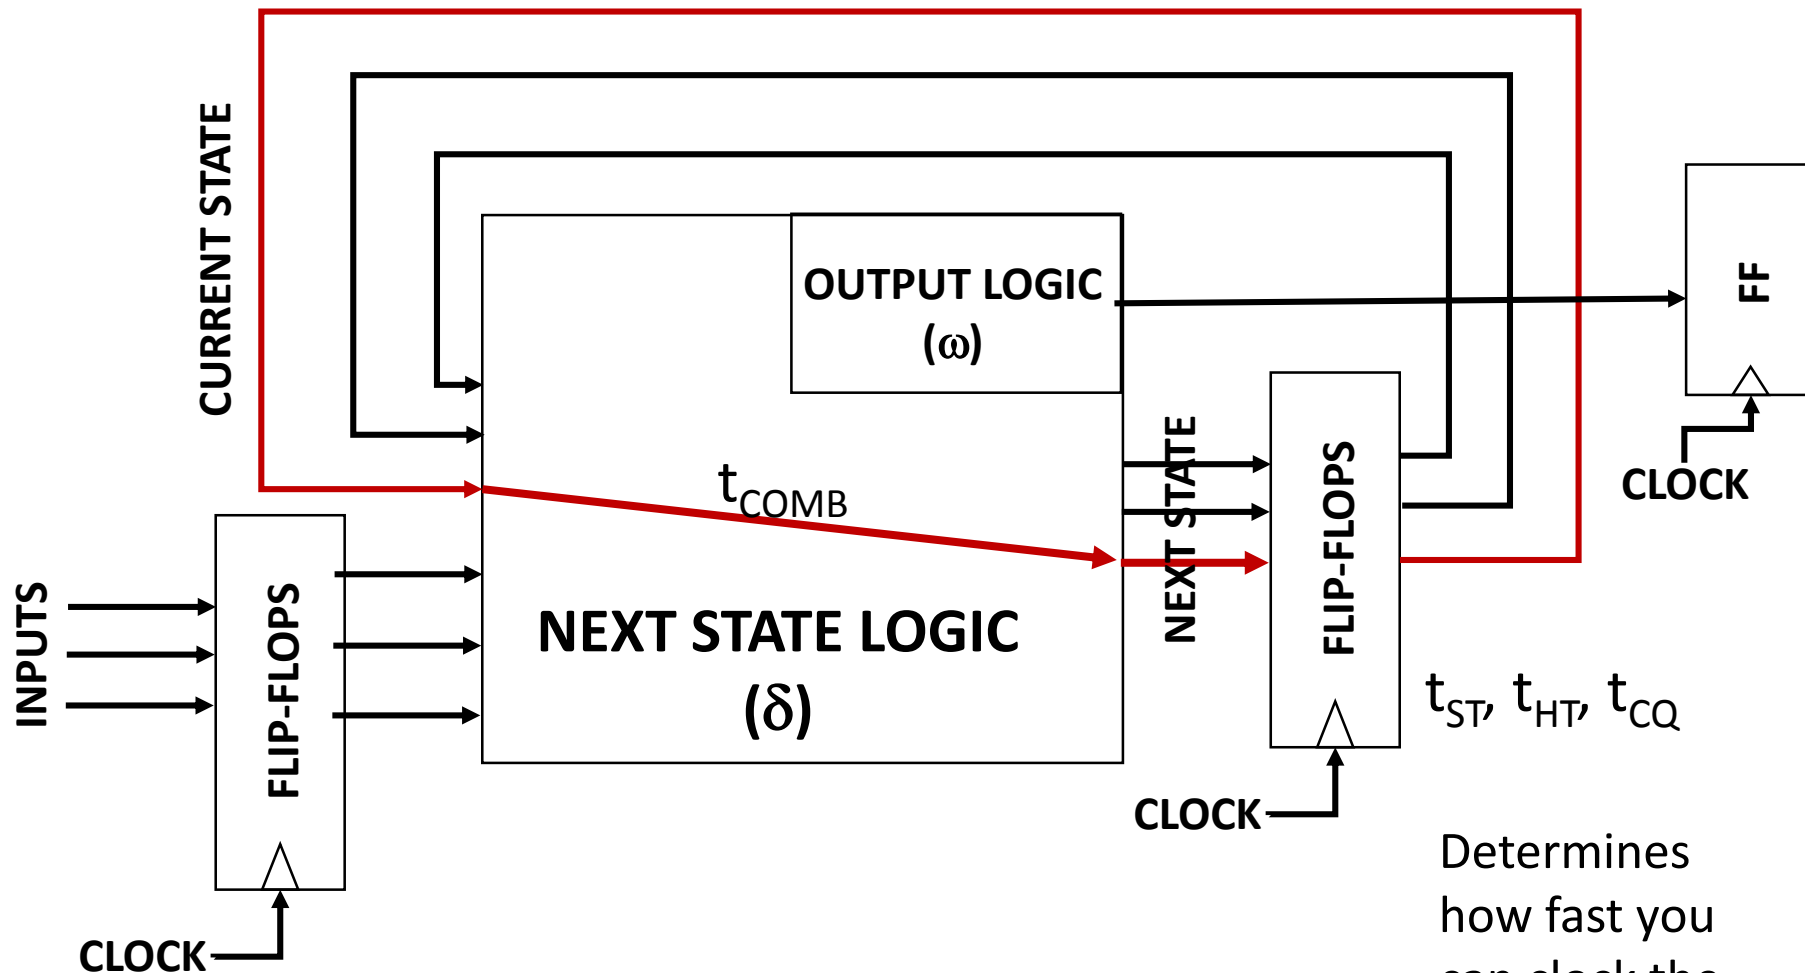
Optimization Metrics

- Three primary metrics that a synthesis tool cares about
 - **Area**: what is the total sum of areas of each gate
 - **Delay**: what is the delay from the input to the output (more on this very shortly)
 - **Power**: what is the circuit power consumption (full lecture dedicated to power consumption later in the term)
- Let's look at delay now
 - But first a word on FSM timing



Flip-flops have three relevant timing parameters:

- **Setup time (ST)**: inputs of FF should be ready ST ns before positive edge
- **Hold time (HT)**: inputs should remain stable for HT ns after pos edge
- **Clock-q (CQ)** propagation delay: outputs ready CQ ns after pos edge



Determines
how fast you
can clock the
circuit


Setup time constraint: $t_{\text{CQ}} + t_{\text{COMB}} + t_{\text{ST}} < T_{\text{CLK}}$

Hold time constraint: $t_{\text{CQ}} + t_{\text{COMB}} > T_{\text{HT}}$

Automatically handled by
most synthesis tools using
buffers

Determining T_{COMB}

- T_{COMB} is the worst-case (largest) delay from an input to an output



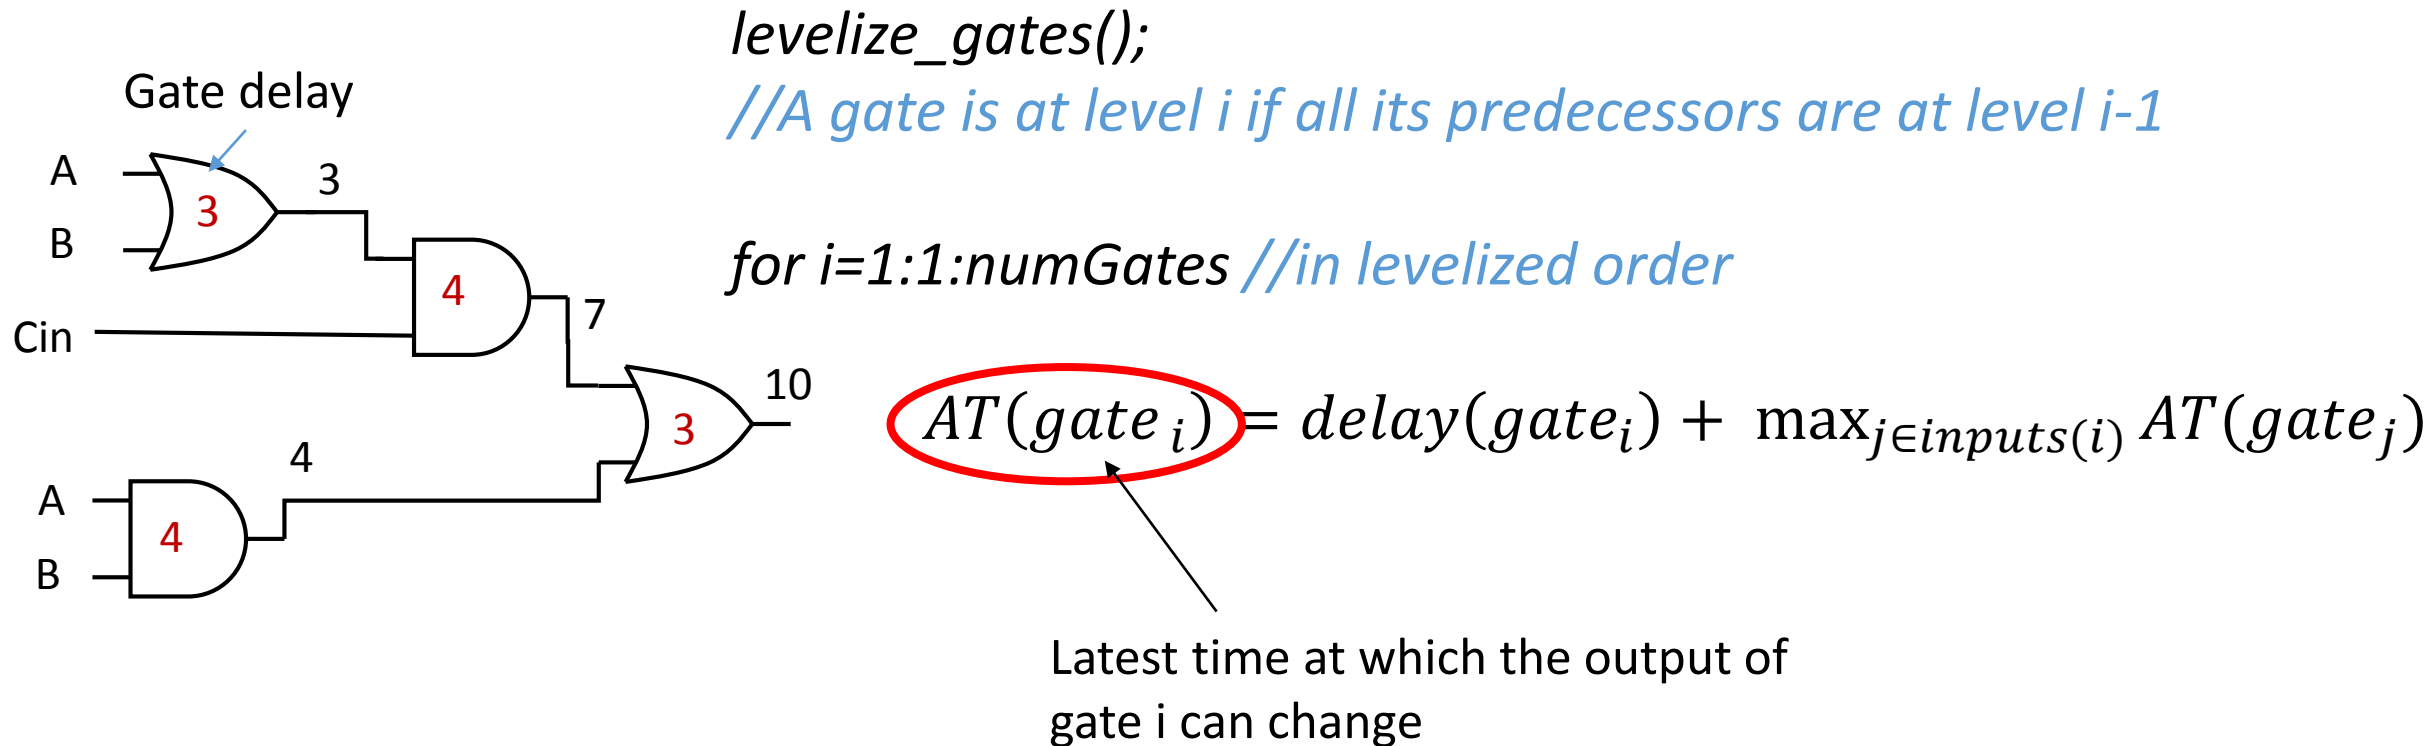
How many input-output paths are there in this netlist?

Determine T_{COMB}

In general, how many paths can there be for a block of combinational logic with N gates?

Static Timing Analysis

- Polynomial time algorithm to determine T_{COMB}



FSM Synthesis

- Simple up-counter

```
module counter (input clk, input reset,  
                input [31:0] step, output reg [31:0] count);  
  
always @ (posedge clk, posedge reset) begin  
    if (reset)  
        count <= 0;  
    else  
        count <= count + step;  
    end  
endmodule
```

- run.tcl script specifies a clock period constraint to RC

```
set myClk clk  
set myPeriod_ps 1000
```


In Class Demo

Target Clock Period	Slack (positive/negative)	Area
10000 ps		
1000 ps		